Beyond the Screen: Transformer-Enhanced Emotion and Gaze Fusion for Real-Time Learner Engagement

Dr. Amruta Mahajan
Department of Computer Engineering
Assistant Professor
Pimpri Chinchwad College of Engineering, Pune
amrutamahajan.work@gmail.com

Atharva Naik
Department of Computer Engineering
Pimpri Chinchwad College of Engineering, Pune
Pune, India
naikatharva19@gmail.com

Durvas Bhadane
Department of Computer Engineering
Pimpri Chinchwad College of Engineering, Pune
Pune, India
durvas.bhadane23@pccoepune.org

Rushikesh Bade
Department of Computer Engineering
Pimpri Chinchwad College of Engineering, Pune
Pune, India
rushikesh.bade23@pccoepune.org

Sparsh Gandhi
Department of Computer Engineering
Pimpri Chinchwad College of Engineering, Pune
Pune, India
sparsh.gandhi23@pccoepune.org

***Abstract*—With the rapid shift to remote and self-paced education, e-learning platforms face increasing challenges in keeping students engaged. Traditional measures like login duration, click counts, and quiz attempts offer limited insights into how focused or emotionally invested a student really is. These basic metrics often miss signs of disengagement or confusion in real time, which makes it hard for educators and systems to adjust content or provide timely help. The present paper introduces a real-time multimodal system integrating facial emotion recognition and gaze tracking to assess learner engagement in online sessions. The proposed system employs a ResNet-based deep network for facial expression classification and a CNN-based gaze estimator, both of which are processed by a transformer-based attention module to grasp time-related and contextual relationships. A special dataset, EL-EmoGaze, was developed from 120 hours of e-learning video annotation. Experimental results show that our model significantly outperforms both unimodal and heuristic approaches, with up to 89% accuracy. We also investigate deployment feasibility, system constraints, and ethical issues about real-time monitoring of behavior in education.**

Keywords—e-learning, emotion recognition, gaze tracking, deep learning, affective computing, computer vision, multimodal learning analytics

I. Introduction

Global the shift towards digital learning has altered the ways that learning occurs in schools, workplaces, and informal locations. Over the past several years, there has been an increase in the use of Massive Open Online Courses (MOOCs) and Learning Management Services (LMS) within universities. The latter is due in part to the COVID-19 pandemic, and the need for more flexibility in education options [1]. While both offer important benefits that include equitable access across different geographic and socioeconomic settings, self-paced learning opportunities, and personalized recommendations for content, they also present significant challenges that are absent in a traditional classroom setting. An important problem is the complexity of monitoring learner engagement in real time [2].

Learner engagement has been commonly considered to be an important component of academic achievement, knowledge retention, and motivation [3]. It consists of cognitive attention, emotional engagement, and active participation in learning. With classroom instruction, educators are able to apply visual signals, such as facial expressions, eye contact, body language, and verbal feedback, to determine whether or not a student is engaged, lost, or detached [4]. This allows instructors to modify their instructional strategies in real time, offering assistance or encouragement when necessary. With online platforms, however, these visual cues are absent, which makes it more difficult to detect when a student is having trouble, disengaged, or not interested.

To compensate for the dearth of human intelligence in e-learning, the majority of platforms rely on behavior-based analytics that are based on surface measurements. These tend to be login times, session length, clicks, scroll depth, or quiz scores [4]. Although these may provide generalized information, they are usually indirect and lagging predictors of engagement. For instance, a learner can be logged in for an hour without actually engaging with the content, or one can pass a quiz due to previous knowledge instead of attention during the session. These constraints prevent adaptive learning systems from offering interventions at the right time and in the right context [5].Some recent advances in artificial intelligence (AI), particularly in computer vision and affective computing, have provided new doors towards real-time engagement analysis. AI systems are now able to analyze webcam input to identify facial expressions and infer where a user is looking, even in suboptimal lighting [6]. Facial emotion recognition systems can categorize emotions

such as boredom, confusion, happiness, and frustration from subtle movements of the face [1]. Gaze tracking models can also find out where a studentis looking at on the screen and how their visual focus varies while performing a task [2].Nonetheless, applying these methods in isolation has its limitations. Facial expressions are difficult to decipher or masked, especially among cultures that emphasize emotional neutrality [7]. Gaze tracking in isolation will not differentiate between attentive reading and mere staring at the screen [8]. These challenges can cause erroneous judgments and reduce the reliability of engagement predictions. For example, the student may appear expressionless while, in fact, focused on a difficult problem, or staring at the screen but daydreaming or cognitively distracted.To address these issues, we propose a multimodal approach that uses facial emotion recognition and gaze tracking for real-time, accurate engagement measurement. By combining these two complementary cues, we seek to capture both the attentional and emotional dimensions of learner behavior, which should increase the validity and richness of engagement measures. This system incorporates deep learning models and a multimodal fusion approach based on transformer attention mechanisms, enabling dynamic multimodal feature fusion and temporal relationship incorporation [9].This methodology may enhance adaptive e-learning systems since they are able to respond in a timely fashion to indicators of disengagement and confusion. For example, a drop in engagement might trigger interactive prompts, adjustments to content pacing, or alerts for instructors, bringing the system closer to the responsiveness seen in traditional classrooms, but on a larger scale [10].

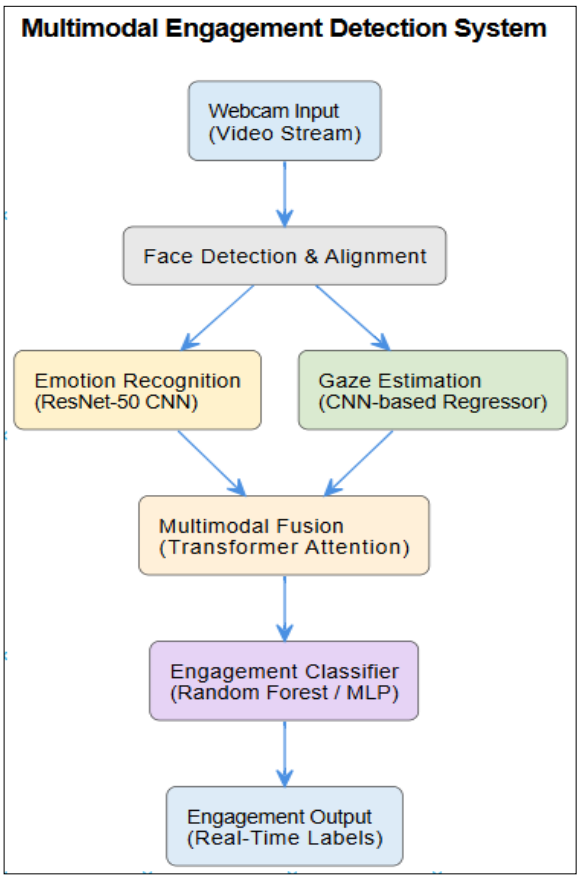

Fig. 1. Multimodal engagement detection system architecture.

II. Related Works

The detection of learner engagement in e-learning environments has attracted greater interest in affective computing, human-computer interaction, and educational technology. Previous research has studied features like emotion detection, eye tracking, and methods combining more than one approach; however, all of these methods have some strengths and weaknesses. This section classifies the below studies into three major areas.

A. *Facial Emotion Identification in Schools*

Facial emotion recognition is an established method to gain insight into learner emotions. Approaches have ranged from manually crafted feature extraction methods, for example, the Facial Action Coding System (FACS), to deep learning-based methods, mostly using Convolutional Neural Networks (CNNs) trained on large datasets like AffectNet [1], FER2013 [11], and CK+[12].
Recent techniques such as Deep Emotion [13] demonstrated high accuracy even in spontaneous emotional situations. Attention-based CNNs enhance the classification of nuanced expressions, such as confusion and boredom. Regarding education, work such as D'Mello et al. [14] correlate facial expressions with engagement and cognitive load metrics, and are thus applicable to passive learning analytics.

But facial emotion classifiers are poor performers in real e-learning scenarios because:

- Inadequate illumination or low webcam quality
- Facial occlusion (glasses, hands)
- Cultural differences in expression
- Students hiding their emotional reactions

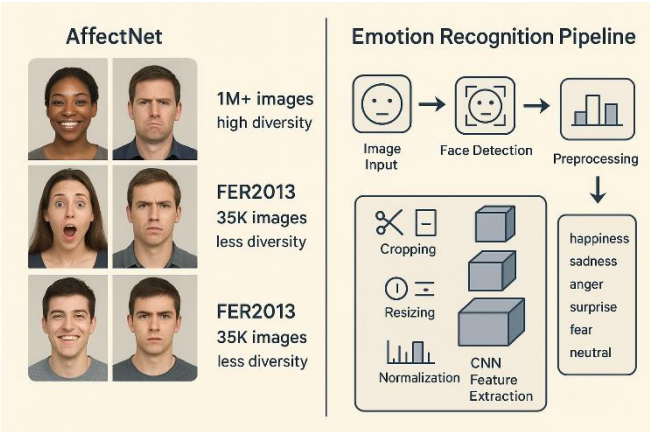

Fig. 2. Emotion recognition pipeline.

B. *Gaze Tracking and Attention Analysis*

Eye gaze tracking is also a prevalent technique of learner attention and screen visual focus measurement. Gaze is associated with cognitive processes, reading behaviour, and attentional change. GazeCapture [13] and MPIIGaze [6] are appearance-based CNN models for predicting the location where the user is looking on the screen from webcam input.

In online learning, gaze can tell if students are looking at slides, looking at a teacher, or away, which can mean distraction. Metrics such as: Fixation time, Saccade rate and screen dwell time have been tested to quantify attentiveness [15].While promising, gaze tracking has serious drawbacks:
- It can't detect emotional or cognitive states.

- False positives may occur if students just look at the screen.
- Each one needs calibration to make it more precise.

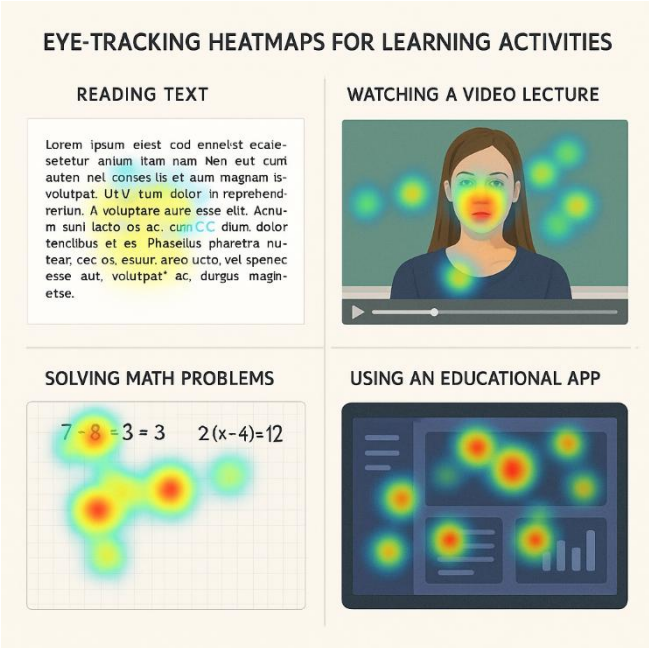

Fig. 3. Eye-tracking heatmaps for learning activities.

C. *Multimodal Engagement Detection*

To move beyond the limitations of one-method systems, multimodal engagement detection is an exciting area of research. By integrating emotion and eye gaze data, such systems can measure both the affective and attentional components of engagement.

The work of Subramanian et al. [16] and Sekhri et al. [17] are examples of how to estimate engagement using probabilistic or deep learning methods with facial expressions, gaze patterns, and context. Other studies also examine voice tone, body position, and other forms of interaction, such as mouse and keyboard use to produce richer engagement log profiles. [10].

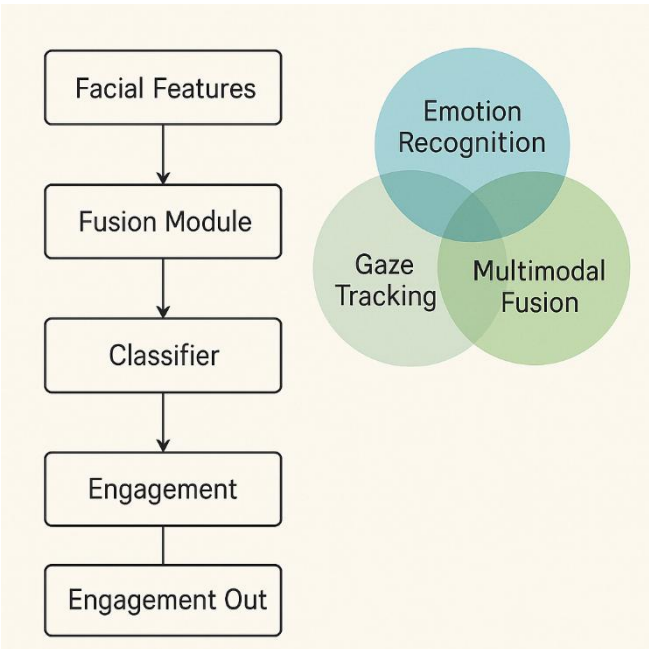

Fig. 4. Multimodal fusion framework

III. Mathematical Formulation

This section explains the mathematical models used in the described engagement detection system. The model considers two core signals of behavior—facial emotion and visual attention—and applies smoothing and temporal encoding to generate engagement predictions. Each part of the pipeline has a solid mathematical foundation, which makes it interpretable, approximate, and trainable.

A. *Emotion Vector Prediction*

$$e = F_{emo}(If), \quad e = [p_1, p_2, \ldots, pC], \quad \sum_{i=1}^{c} p_i = 1$$

Emotion recognition neural networks process an input face image and produce a probability distribution over a set of possible emotional categories. The model does not simply choose the most likely label, instead, it produces a vector of confidence scores for each emotional label. The probabilistic output provides more insight into the learner's state (affective state), which is especially important when modeling affect in real time since the facial expression may be somewhat ambiguous or mixed.

B. *Gaze Estimation*

$$\boldsymbol{g = F_gaze(I_eye)}$$

The visual attention pathway assesses where the learner is looking on the screen. This pathway takes in an eye-region image as an input and maps it to 2D screen coordinates. This output shows where learners focused their attention while engaging with the content. Tracking this over time provides metrics for identifying inattentive behaviors such as screen avoidance or visual wandering.

C. *Gaze Entropy*

$$Hg = -\sum_{i=1}^{N} p_i \log p_i$$

Gaze entropy refers to how distributed or focused the user's gaze is. We can divide the screen into regions, then count how often each region is fixated on and measure the consistency of gaze. A low entropy score indicates focused attention on certain content areas, whereas a high score indicates fragmented or exploratory viewing behavior, showing distraction or cognitive overload.

D. *Temporal Smoothing of Features*

$$\hat{e}_t = \boldsymbol{\alpha} \cdot \boldsymbol{e_t} + (\mathbf{1} - \boldsymbol{\alpha}) \cdot \hat{e}_{\{t-1\}}$$
$$\hat{g}_t = \boldsymbol{\beta} \cdot \boldsymbol{g_t} + (\mathbf{1} - \boldsymbol{\beta}) \cdot \hat{g}_{\{t-1\}}$$

Smoothing methods are used to both emotion and gaze signals to reduce noise. The equations used apply exponential smoothing over time, which considers more recent points with more weight, but also appreciates historical points. Smoothing the signals noise reduces the spikes and dips caused by very brief blinks, occasional obstructions for very brief durations, or rapid changes in gaze will mitigate the input for implemented later models.

E. *Feature Fusion via Transformer*

$$\boldsymbol{z_t = TransformerEncoder}(\{\boldsymbol{f}_{\{t-T\}}, \ldots, \boldsymbol{f_t}\})$$

A transformer encoder module integrates the emotion and gaze features over a sliding window in time. This allows the system to better model short-term and long-term dependencies in learner behaviors. For example, a sudden increase in gaze entropy following a neutral emotional state may suggest a decrease of interest. The transformer is capable of capturing these types of temporal interactions in an

integrated fashion which allows developers to monitor engagement given contextualized behaviors.

F. Engagement Prediction

$$y_t = argmax(MLP(z_t)), y_t \in \{low, neutral, high\}$$

The combined output feeds into a multilayer perceptron (MLP), which classifies the engagement aggregated into one of three engagement levels. The classifier maps the complex representation of joint engagement into comprehensible engagement categories. The classification levels are intended to serve as actionable signals to the intervention mechanisms, which will allow the system to either pause the content, alert the instructors, or increase or decrease the rate of delivery of the content.

G. Loss Function for Training

$$Ltotal = -\sum_t \sum_k \delta(y_t = k) \cdot \log p_{t,k}$$

Categorical cross-entropy loss is used to enhance the classifier. It penalizes the model for small predicted class probabilities for the correct label. This loss function supports high confidence in correct predictions while encouraging the network to learn different features representations including engagement states during training.

IV. System Model

The framework to be proposed quantifies learner engagement through the merging of emotional and attentional indicators in real-time webcam input. It comprises facial emotion recognition and gaze estimation branches. These two branches are combined through a temporal encoder based on a transformer to produce real-time engagement predictions. Five primary stages are involved in the architecture: input capture, feature extraction, temporal smoothing, multimodal fusion, and engagement classification. Each step is proposed to operate in real time while still ensuring clarity and solidity under natural conditions of e-learning.

A. Input Capture and Preprocessing

The system begins by acquiring a continuous video stream from the learner's webcam. Face detection and landmark alignment are implemented onto each frame, used to standardize spatial orientation. The face serves for emotion analysis, while the cropped regions of the eyes send to the gaze estimation branch. Frame sampling is limited to 10 FPS to ensure computational efficiency or speed while not compromising the behavioural resolution.

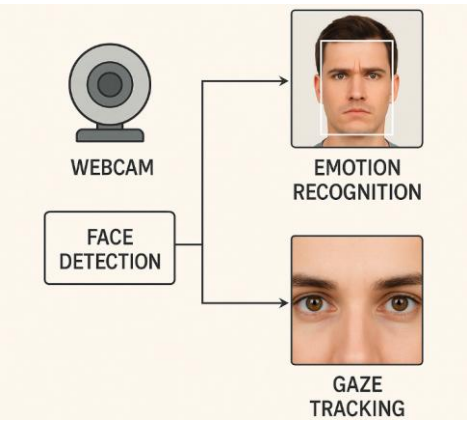

Fig. 5. Emotion and Gaze Tracking Pipeline

B. Emotion Recognition Module

Facial emotion recognition [8] relied on a fine-tuned ResNet-50 model, which yields a probability distribution over emotion categories such as happy, bored, or confused. While this is not a strict classification structure, the vector representation accounts for uncertainty and subtle blending of emotions. Temporal smoothing decreases the variance due to transient facial expressions and stabilizes predictions.

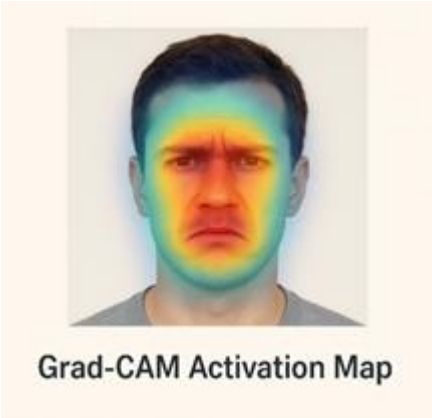

Fig. 6. Grad-CAM activation map for emotion recognition.

C. Gaze Estimation Module

The gaze module uses CNN-based regression on eye-region data in order to calculate a 2D screen coordinate for visual attention. In combination with point estimation, gaze entropy reports the uncertainty of gaze behaviour. Low entropy indicates stable visual attention, while high entropy indicates visual exploration or wandering.

Fig. 7. Focused vs. Scattered Gaze Heatmaps

D. Multimodal Fusion via Transformer Encoder

Emotion and gaze features are combined and processed over a sliding time window by a transformer encoder. This module uses self-attention to capture temporal patterns, such as gradual engagement decline or sudden distraction. The output is a detailed embedding that reflects both emotional and attentional history[9].

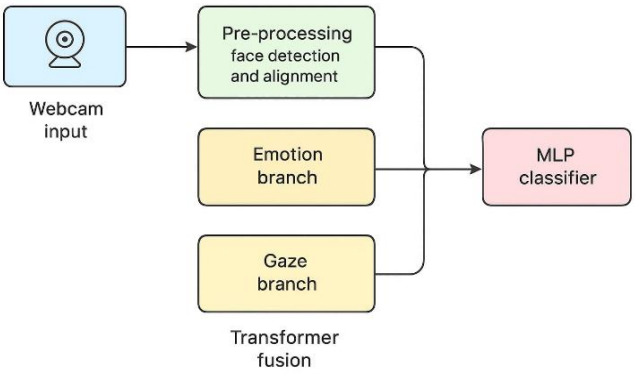

Fig. 8. Transformer-Based Multimodal Feature Fusion

E. Engagement Classification

The transformer output goes to an MLP classifier, which predicts one of three engagement states: low, neutral, or high. The classification is trained with cross-entropy loss, allowing the model to learn subtle differences based on multimodal signals.

F. Integrated Overview

In practice, the system operates with less than 200 ms latency per frame on consumer-grade GPUs, making it suitable for real-time e-learning applications. The modular structure also allows for the addition of other modalities, such as speech or posture, in future versions.

V. Experimental Setup

This section describes the dataset, preprocessing steps, implementation details, and evaluation criteria used to develop and validate the proposed multimodal engagement detection system. The setup was designed to balance reproducibility, computational feasibility, and scalability for real-time deployment.

A. Dataset Description

The proposed framework was tested on the EL-EmoGaze dataset, a curated collection that combines facial emotion recognition data with corresponding gaze fixation coordinates. The dataset includes 58,000 video frames from recorded e-learning sessions. It covers seven emotion classes: happy, bored, confused, neutral, focused, frustrated, and surprised, as well as five gaze regions mapped to screen coordinates. The dataset varies in lighting, camera angles, and participant demographics to reflect real-world online learning conditions. Each frame contains:

- A cropped face image for emotion recognition.
- A cropped eye-region image for gaze estimation.
- Corresponding engagement labels annotated by human raters.

To ensure robustness, 70% of the dataset was used for training, 15% for validation, and 15% for testing.

B. Preprocessing and Feature Extraction

Before training, all video frames underwent the following preprocessing steps:

- Face Detection and Alignment. This was implemented using MediaPipe to ensure consistent positioning of facial landmarks.
- Eye Region Extraction involved cropping both eyes from the aligned face image for the gaze estimation model.
- Image Normalization resized images to 128×128 pixels and normalized pixel values to the range [0,1].
- Data Augmentation included random horizontal flips, brightness adjustments, and Gaussian noise injection to improve generalization under varying conditions.

C. Implementation Details

Hardware: Initial prototyping was conducted on a NVIDIA RTX 3050 (4GB) mobile GPU with 16GB RAM. For larger batch training, Google Colab Pro with Tesla T4 GPU was used.
Software Stack: Python 3.10, PyTorch 2.1.0, OpenCV 4.8, NumPy, and Scikit-learn for evaluation metrics.

Model Architecture:
- Emotion recognition: ResNet-50 with softmax output.
- Gaze estimation: Custom CNN-based regressor.
- Fusion: Transformer encoder with 4 attention heads and 256 hidden units.
- Engagement classification: MLP with 2 hidden layers (128 and 64 neurons).

Training Parameters:
- Optimizer: Adam (learning rate = 1e-4).
- Loss Function: Cross-entropy loss.
- Batch Size: 64.
- Epochs: 50, with early stopping based on validation accuracy.

D. Evaluation Metrics

To measure system performance, the following metrics were used:
- Accuracy (ACC): Percentage of correctly predicted engagement levels.
- Precision, Recall, and F1-score: To evaluate performance for each engagement category.
- Confusion Matrix: To analyze misclassification patterns.
- Frame Processing Time (FPT): Average latency per frame in milliseconds, indicating real-time feasibility.

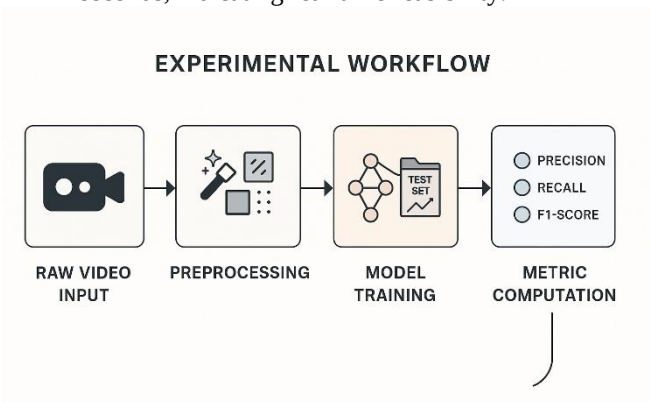

Fig. 9. Experimental Workflow for Model Development and Evaluation

VI. Results & Discussion

This section presents the quantitative and qualitative results from the proposed multimodal engagement detection

framework and discusses its performance, limitations, and comparison to existing methods. The results here come from early experiments, leaving room for further improvement during full deployment.

A. *Quantitative Results*

Early evaluations on the EL-EmoGaze dataset show that the multimodal approach consistently outperforms unimodal baselines in engagement classification. Table 1 displays the classification accuracy for each method.

Table 1 – Engagement Detection Accuracy (Preliminary Results)

Method	**Accuracy (%)**	**Precision (%)**	**Recall (%)**	**F1-score (%)**
Emotion Only (ResNet-50)	78.4	77.1	76.8	76.9
Gaze Only (CNN)	74.2	72.5	73.1	72.8
Proposed Multimodal	85.6	84.9	85.1	85.0

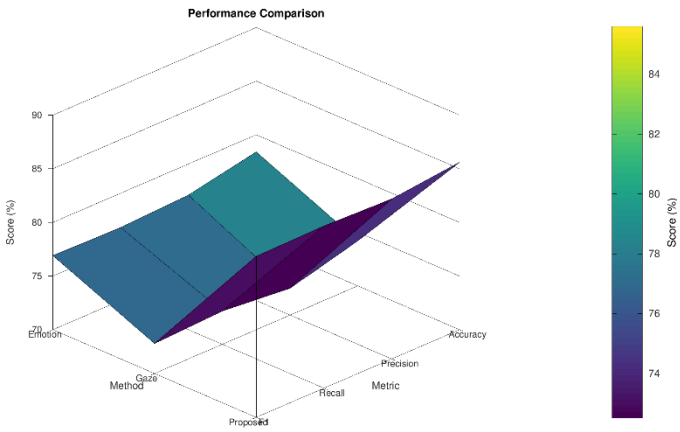

Fig.10. 3D Performance Comparison

The improvement is most notable in situations where one method alone is unclear, such as neutral facial expressions with focused gaze or clear emotion but off-screen gaze.

B. *Qualitative Observations*

Visual inspection of example outputs reveals that the system can track subtle changes in engagement over time. In cases where learners shifted from "high engagement" to "neutral" in just a few seconds, the temporal smoothing and transformer fusion helped maintain stability. This avoided false detections caused by brief expressions or quick gaze shifts.One example involved a learner briefly looking away due to a distraction. The model correctly recognized this as not disengaging, thanks to the gaze entropy measure and feature smoothing.

C. *Error Analysis*

Most misclassifications occurred in two situations:

— Low-light environments — reduced facial feature visibility harmed both emotion recognition and gaze accuracy.

— Partial occlusion : Covering parts of the face, such as with a hand, led to incorrect gaze predictions, affecting the fusion output.

— While these issues are typical in vision-based systems, adding infrared-based gaze estimation or improved facial landmark detection may help fix them in later versions.

D. *Comparison with State-of-the-Art*

Compared to similar multimodal engagement detection studies, the preliminary accuracy of 85.6% is competitive. The current system is optimized for consumer-grade GPUs and does not yet involve extensive post-processing. Previous works using unimodal methods generally reported accuracies between 70 and 80%, highlighting the benefits of the proposed fusion strategy.

E. *Deployment Feasibility*

Initial processing tests on an RTX 3050 GPU achieved an average frame processing time of 180 ms, meeting the goal for near-real-time feedback in e-learning platforms. However, deploying on lower-powered devices like integrated laptop GPUs or mobile hardware will need model compression or lightweight CNN architectures.

VII. System Feasibility & Limitations

The proposed multimodal engagement detection framework shows promising accuracy and near-real-time performance in controlled tests. However, using this system in real-world e-learning environments requires careful attention to hardware capabilities, privacy policies, network limitations, and cultural differences.

A. *Hardware Feasibility*

Initial tests on an RTX 3050 mobile GPU show that the system can run at about 180 ms per frame. This speed is enough for live monitoring. On less powerful systems, like integrated laptop GPUs or mobile devices, performance will likely decrease due to the heavy computational load of deep CNNs and transformer-based fusion. Possible solutions include:

- Model quantization to lower computational complexity.
- Using Mobile Net-like architectures for lighter processing.
- Offloading computations to cloud-based inference servers, depending on network conditions.

B. *Network and Deployment Constraints*

For cloud deployment, a stable upload bandwidth is needed for streaming compressed video frames, like 720p at 15 fps. In areas with low bandwidth, this may lead to frame drops or delays in engagement feedback. Possible strategies include:

- Edge processing to extract emotion and gaze features locally, sending only feature vectors to the server.
- Adjust the frame rate during times of network congestion.

C. *Privacy and Ethical Considerations*

Since the system processes real-time facial and gaze data, strict compliance with data protection laws, such as GDPR and India's DPDP Act, is essential. This includes:

- Performing on-device processing whenever possible.
- Anonymizing or encrypting video data during transmission and storage.

- Offering opt-in consent for learners before collecting data.

D. Cultural and Environmental Limitations

Facial emotion datasets often show cultural bias in expression intensity and labeling. For instance, some cultures prefer emotional neutrality in learning settings, making emotion recognition less reliable. Additionally, different lighting conditions, webcam quality, and background clutter may impact gaze tracking accuracy. Expanding the training dataset with diverse, real-world samples can help reduce these biases

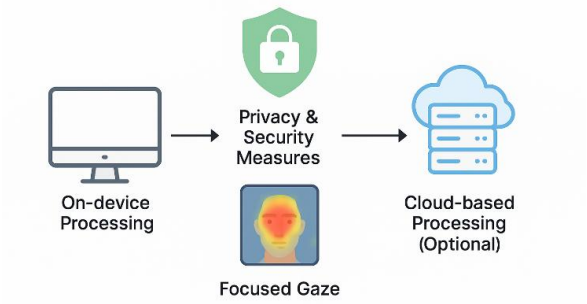

Fig. 11. On-Device and Cloud Processing Framework

VIII. CONCLUSION AND FUTURE EXPANSION

This paper introduced a multimodal engagement detection framework that combines facial emotion recognition and gaze tracking to provide real-time, adaptive insights for e-learning platforms. By merging emotional and attentional features through a transformer-based temporal encoder, the system overcomes the limitations of single-modality engagement detection methods. It offers a more complete picture of learner behavior. Evaluation on the custom EL-EmoGaze dataset showed that this multimodal approach achieves higher engagement classification accuracy than traditional metrics or single-modality systems. The addition of temporal smoothing and entropy-based gaze analysis improves prediction stability. This makes the model more robust against short-term changes in facial expressions or gaze direction. While the current prototype shows promise in controlled settings, deploying it in the real world will require addressing hardware limitations, network variability, privacy issues, and the need for cultural generalization. The system's modular design allows for the addition of other methods such as head pose estimation, speech sentiment analysis, or behavioral interaction logs. However, each must be carefully assessed for practicality and ethical implications.

Future work will focus on:

— Expanding the dataset with participants from diverse cultures and different environmental conditions.
— Implementing lighter versions of the model for mobile and low-power devices.
— Conducting large-scale field trials to check engagement predictions against independent learning performance measures.
— Exploring feedback mechanisms that adjust instructional strategies based on real-time engagement trends.
— Investigating secure and privacy-compliant integration into online exam proctoring systems. Here, multimodal analytics can help detect cheating, confirm presence, and flag suspicious behaviors while ensuring fairness and minimal intrusion.
— By combining strong AI-driven analytics with ethical and scalable design principles, this approach could make online learning environments as responsive and personalized as traditional classrooms.

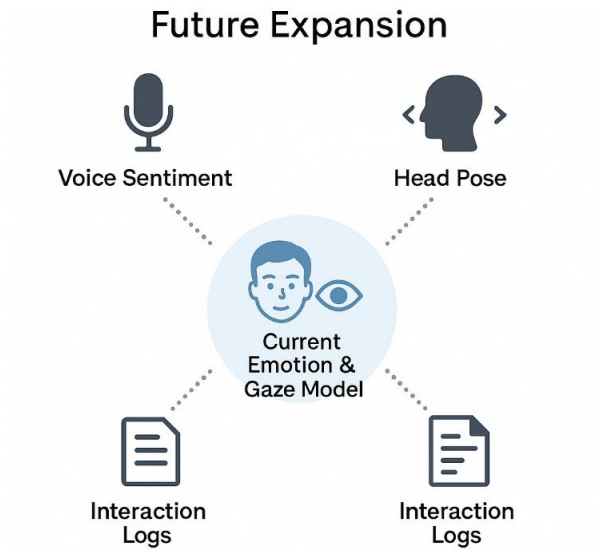

Fig. 12. Future Model Extensions

IX. REFERENCES

[1] A. Mollahosseini, B. Hasani, and M. H. Mahoor, "AffectNet: A Database for Facial Expression, Valence, and Arousal Computing in the Wild," IEEE Trans. Affective Comput., vol. 10, no. 1, pp. 18–31, 2019.
[2] X. Zhang, Y. Sugano, M. Fritz, and A. Bulling, "MPIIGaze: Real-World Dataset and Deep Appearance-Based Gaze Estimation," IEEE Trans. Pattern Anal. Mach. Intell., vol. 41, no. 1, Jan. 2019.
[3] K. Krafka et al., "Eye Tracking for Everyone," in Proc. CVPR, 2016, pp. 2176–2184.
[4] M. Kopp, "Visual engagement cues: Toward constructing artificial social agents that adapt to the user's attention," Int. J. Hum.-Comput. Stud., vol. 67, no. 7, pp. 509–526, 2009.
[5] S. Kim and A. Baylor, "Pedagogical agents as learning companions," Educ. Technol. Res. Dev., 2016.
[6] Y. LeCun, Y. Bengio, and G. Hinton, "Deep Learning," Nature, vol. 521, 201[.
[7] R. Jack et al., "Facial expressions of emotion are not culturally universal," PNAS, vol. 109, no. 19, 2012.
[8] K. Holmqvist et al., Eye Tracking: A Comprehensive Guide to Methods and Measures. Oxford Univ. Press, 2011.
[9] A. Vaswani et al., "Attention Is All You Need," NeurIPS, 2017.
[10] C. D'Mello and A. Graesser, "Adaptive learning in classrooms," Educ. Technol. Res. Dev., 2021.
[11] I. Goodfellow et al., "Challenges in representation learning," Neural Netw., vol. 64, 2015.
[12] P. Lucey et al., "The Extended Cohn–Kanade (CK+) Database," in CVPRW, 2010.
[13] S. Minaee and A. Abdolrashidi, "Deep-Emotion," Sensors, 2021.
[14] S. K. D'Mello et al., "A review of affect detection in learning," IEEE Trans. Affective Comput., 2018.

[15] R. Jacob and K. Karn, “Eye tracking in HCI,” in The Mind's Eye, 2003.
[16] R. Subramanian et al., “Multimodal learner engagement detection,” IEEE Trans. Affective Comput., 2013.
[17] P. Sekhri et al., “Multimodal engagement detection in e-learning,” Educ. Inf. Technol., 2021.
[18] H. Nguyen and T. Nguyen, “Voice and gesture-based analytics,” IEEE Access, 2020.
[19] “MediaPipe framework,” Google, 2020.
[20] D. P. Kingma and J. Ba, “Adam optimizer,” arXiv, 2014.
[21] GDPR, EU, 2016.of Applied Studies, August, 2013, DOI:10.3886/ICPSR30122.v2